

# SQL Interview Questions for Experienced Candidates (3+ years)

## Data Types & Miscellaneous Concepts

## Normalization & Schema Design

## Database Design Principles

## Indexing Strategies & Keys

## Indexing Strategies & Keys

### 1. What is a SQL index and what are different types of indexes (clustered, non-clustered, unique, etc.)?

#### Answer:

An **index** in SQL is a database object that improves the speed of data retrieval operations on a table at the cost of additional storage and maintenance overhead. Indexes work like a book's index, allowing the database engine to find rows quickly without scanning the entire table.

- **Clustered Index:** Determines the physical order of data in the table. Each table can have only one clustered index.
- **Non-Clustered Index:** A separate structure from the table data, containing pointers to the actual data rows. A table can have multiple non-clustered indexes.
- **Unique Index:** Ensures that all values in the indexed column(s) are unique.
- **Composite Index:** An index on two or more columns.
- **Full-Text Index:** Used for efficient text searching.
- **Spatial Index:** Used for spatial data types (e.g., geometry, geography).

#### Syntax & Example:

- **Clustered Index:**

```
CREATE CLUSTERED INDEX idx_employee_id ON Employees(EmployeeID);
```

This creates a clustered index on the *EmployeeID* column of the *Employees* table. The data rows will be physically ordered by *EmployeeID*.

- **Non-Clustered Index:**

```
CREATE NONCLUSTERED INDEX idx_employee_lastname ON Employees(LastName);
```

This creates a non-clustered index on the *LastName* column. The index is stored separately from the table data and contains pointers to the actual rows.

- **Unique Index:**

```
CREATE UNIQUE INDEX idx_employee_email ON Employees(Email);
```

This ensures that all values in the *Email* column are unique.

- **Composite Index:**

```
CREATE INDEX idx_employee_dept_salary ON Employees(DepartmentID, Salary);
```

This creates an index on both *DepartmentID* and *Salary*. Useful for queries filtering or sorting by both columns.

**Summary:** Indexes speed up SELECT queries but can slow down data modification operations (INSERT, UPDATE, DELETE).

---

## 2. What is the difference between a heap (no clustered index) and a table with a clustered index, and how can you identify a heap table?

### Answer:

A **heap** is a table without a clustered index. Data is stored in no particular order, and row locations are tracked by row identifiers (RIDs). A table with a **clustered index** stores data rows in the order of the index key.

- **Heap Table:** No clustered index; data is unordered. Identified by querying system catalog views (e.g., *sys.indexes* in SQL Server) and checking for *index\_id = 0*.
- **Clustered Index Table:** Data is physically ordered by the clustered index key; *index\_id = 1*.

### Syntax & Example:

- **Create a heap (no clustered index):**

```
CREATE TABLE HeapTable (  
    ID INT,  
    Name VARCHAR(50)  
);  
-- No clustered index created, so this is a heap.
```

This table is a heap because no clustered index is defined.

- **Add a clustered index:**

```
CREATE CLUSTERED INDEX idx_id ON HeapTable(ID);
```

Now, *HeapTable* is no longer a heap; it is ordered by *ID*.

- **Identify a heap in SQL Server:**

```
SELECT name, index_id
FROM sys.indexes
WHERE object_id = OBJECT_ID('HeapTable');
```

If *index\_id = 0* exists, the table is a heap.

**Summary:** Heaps are faster for bulk inserts but slower for searches; clustered indexes improve search performance.

---

### 3. What is the difference between a PRIMARY KEY and a UNIQUE KEY (or unique index) in SQL?

**Answer:**

Both **PRIMARY KEY** and **UNIQUE KEY** enforce uniqueness on columns, but there are differences:

- **PRIMARY KEY:** Uniquely identifies each row; only one per table; cannot contain NULLs; automatically creates a unique clustered index (if none exists).
- **UNIQUE KEY:** Enforces uniqueness; multiple unique keys allowed per table; columns can contain a single NULL (in most databases); creates a unique non-clustered index by default.

**Syntax & Example:**

- **PRIMARY KEY:**

```
CREATE TABLE Employees (
    EmployeeID INT PRIMARY KEY,
    Email VARCHAR(100) UNIQUE
);
```

*EmployeeID* is the primary key (unique, not null). *Email* is a unique key (can be null in some databases).

- **UNIQUE KEY (explicit):**

```
ALTER TABLE Employees ADD CONSTRAINT uq_emp_phone UNIQUE (PhoneNumber);
```

This adds a unique constraint to *PhoneNumber*.

**Summary:** Use PRIMARY KEY for the main identifier; use UNIQUE KEY for alternate unique constraints.

---

### 4. What are index “forwarding pointers” in a heap table, and how do they affect query performance?

**Answer:**

In a **heap table**, when a row is updated and no longer fits in its original location, it may be moved elsewhere. A **forwarding pointer** is left at the original location, pointing to the new location.

- **Impact:** Causes extra I/O because the database must follow the pointer to find the actual row, slowing down queries.
- **Resolution:** Rebuilding the table or adding a clustered index removes forwarding pointers.

**Example:**

```
-- Update a row in a heap table that causes it to move
UPDATE HeapTable SET Name = REPLICATE('A', 1000) WHERE ID = 1;
-- This may create a forwarding pointer if the row can't fit in its original page.
```

Queries that access this row will incur extra I/O to follow the pointer.

**Summary:** Forwarding pointers degrade performance in heaps with frequent updates.

---

## 5. What is a composite index, and how do you choose the order of columns in it for optimal performance?

### Answer:

A **composite index** is an index on two or more columns. The order of columns matters for query optimization.

- **Column Order:** Place the most selective (most unique) column first, or the column most often used in WHERE or JOIN conditions.
- **Index Usage:** The index is most effective when queries filter on the leading column(s).

### Syntax & Example:

```
CREATE INDEX idx_dept_salary ON Employees(DepartmentID, Salary);
```

This index is useful for queries like *WHERE DepartmentID = ? AND Salary > ?*. If you filter only on *Salary*, the index may not be used efficiently.

**Summary:** Choose column order based on query patterns and selectivity.

---

## 6. When should you use a covering index, and how does it improve the performance of a query?

### Answer:

A **covering index** includes all columns needed by a query (in the index key or as included columns), so the database can satisfy the query using only the index, without accessing the table data.

- **Use Case:** For frequently run queries that select a small set of columns.
- **Performance:** Reduces I/O and improves speed by avoiding lookups in the base table (bookmark lookups).

### Syntax & Example:

```
CREATE INDEX idx_covering ON Employees(DepartmentID) INCLUDE (Salary, FirstName);
```

This index covers queries like *SELECT Salary, FirstName FROM Employees WHERE DepartmentID = ?* because all needed columns are in the index.

**Summary:** Covering indexes are powerful for read-heavy workloads with predictable queries.

---

## 7. How does the existence of an index on a column affect INSERT, UPDATE, and DELETE performance on a table?

### Answer:

Indexes speed up SELECT queries but add overhead to data modification operations.

- **INSERT:** Indexes must be updated for each new row, increasing insert time.
- **UPDATE:** If indexed columns are updated, the index must be modified, adding overhead.
- **DELETE:** Index entries must be removed, which can slow down deletes.

### Example:

```
-- Insert into a table with indexes
INSERT INTO Employees (EmployeeID, FirstName, LastName) VALUES (1, 'Ashish', 'Zope');
-- The database updates all relevant indexes after the insert.
```

More indexes mean more work for each insert, update, or delete operation.

**Summary:** More indexes = faster reads, slower writes. Balance based on workload.

---

## 8. What is index selectivity, and why is it important for query optimization?

### Answer:

**Index selectivity** is the ratio of the number of distinct values in an indexed column to the total number of rows. High selectivity means many unique values; low selectivity means many duplicates.

- **Importance:** High selectivity indexes are more useful for filtering queries, as they reduce the number of rows scanned.
- **Low Selectivity:** Indexes on columns with few unique values (e.g., gender) are less effective.

### Example:

```
-- High selectivity: EmployeeID (unique for each row)
CREATE INDEX idx_employee_id ON Employees(EmployeeID);

-- Low selectivity: Gender (few unique values)
CREATE INDEX idx_gender ON Employees(Gender);
```

The `idx_employee_id` index is highly selective and efficient for lookups. The `idx_gender` index is less useful because many rows share the same value.

**Summary:** Use indexes on columns with high selectivity for best performance.

---

## 9. How many clustered indexes can a table have, and why?

### Answer:

A table can have **only one clustered index** because the data rows can be physically ordered in only one way.

- **Reason:** The clustered index defines the physical storage order of the table.
- **Non-Clustered Indexes:** Multiple non-clustered indexes are allowed.

### Syntax & Example:

```
-- Only one clustered index allowed
CREATE CLUSTERED INDEX idx_emp_id ON Employees(EmployeeID);

-- Multiple non-clustered indexes allowed
CREATE NONCLUSTERED INDEX idx_emp_email ON Employees(Email);
CREATE NONCLUSTERED INDEX idx_emp_phone ON Employees(PhoneNumber);
```

Attempting to create a second clustered index will result in an error.

**Summary:** One clustered index per table; unlimited non-clustered indexes (within system limits).

---

## 10. What is index fragmentation, and how can it be resolved or mitigated in a large database?

### Answer:

**Index fragmentation** occurs when the logical order of index pages does not match the physical order, leading to inefficient I/O and slower queries.

- **Causes:** Frequent INSERT, UPDATE, DELETE operations.
- **Resolution:** Rebuild or reorganize indexes using database maintenance commands (e.g., `ALTER INDEX REBUILD` or `REORGANIZE` in SQL Server).
- **Mitigation:** Schedule regular index maintenance, monitor fragmentation levels.

### Syntax & Example:

```
-- Rebuild an index (removes fragmentation)
ALTER INDEX idx_emp_id ON Employees REBUILD;

-- Reorganize an index (less intensive)
ALTER INDEX idx_emp_id ON Employees REORGANIZE;
```

Use these commands regularly to keep indexes efficient, especially in large, busy databases.

**Summary:** Regular index maintenance is essential for optimal performance in large databases.

---

## Mastering SQL Joins

---

### 1. What are the different types of SQL joins (`INNER JOIN`, `LEFT JOIN`, `RIGHT JOIN`, `FULL JOIN`, `CROSS JOIN`, etc.) and when would you use each?

#### Answer:

SQL joins are used to combine rows from two or more tables based on related columns. The main types are:

- **INNER JOIN:** Returns only rows with matching values in both tables. Use when you need records present in both tables.
- **LEFT JOIN (LEFT OUTER JOIN):** Returns all rows from the left table and matched rows from the right table. Unmatched rows from the right table return NULLs. Use when you want all records from the left table, regardless of matches.
- **RIGHT JOIN (RIGHT OUTER JOIN):** Returns all rows from the right table and matched rows from the left table. Unmatched rows from the left table return NULLs. Use when you want all records from the right table.
- **FULL JOIN (FULL OUTER JOIN):** Returns all rows when there is a match in either table. Unmatched rows from either side return NULLs. Use when you want all records from both tables.
- **CROSS JOIN:** Returns the Cartesian product of both tables (every row of the first table joined with every row of the second). Use rarely, typically for generating combinations.

#### Example:

##### INNER JOIN:

```
SELECT a.*, b.*
FROM TableA a
INNER JOIN TableB b ON a.id = b.a_id;
```

##### LEFT JOIN:

```
SELECT a.*, b.*  
FROM TableA a  
LEFT JOIN TableB b ON a.id = b.a_id;
```

**RIGHT JOIN:**

```
SELECT a.*, b.*  
FROM TableA a  
RIGHT JOIN TableB b ON a.id = b.a_id;
```

**FULL OUTER JOIN:**

```
SELECT a.*, b.*  
FROM TableA a  
FULL OUTER JOIN TableB b ON a.id = b.a_id;
```

**CROSS JOIN:**

```
SELECT a.*, b.*  
FROM TableA a  
CROSS JOIN TableB b;
```

**Explanation:**

- Use **INNER JOIN** when you only want rows with matches in both tables.
- Use **LEFT JOIN** to get all rows from the left table, even if there are no matches in the right.
- Use **RIGHT JOIN** to get all rows from the right table, even if there are no matches in the left.
- Use **FULL OUTER JOIN** to get all rows from both tables, with NULLs where there are no matches.
- Use **CROSS JOIN** to get every combination of rows from both tables (rarely used in practice).

---

**2. What is the difference between a **CROSS JOIN** and a **FULL OUTER JOIN**?****Answer:**

**CROSS JOIN** and **FULL OUTER JOIN** are both used to combine rows from two tables, but they operate very differently:

**Join Type Comparison**

Feature	CROSS JOIN	FULL OUTER JOIN
Purpose	Returns the Cartesian product of both tables (all possible combinations of rows).	Returns all rows from both tables, matching rows where possible, and filling with NULLs where there is no match.
Join Condition	No join condition is used.	Join condition is required (typically ON clause).
Result Size	Number of rows = rows in TableA × rows in TableB.	Number of rows = all matched rows + unmatched rows from both tables.

Feature	CROSS JOIN	FULL OUTER JOIN
Typical Use Case	Generating all possible combinations (e.g., scheduling, permutations).	Combining all data from both tables, showing matches and non-matches.

**Example:****CROSS JOIN:**

```
SELECT a.*, b.*
FROM TableA a
CROSS JOIN TableB b;
```

**FULL OUTER JOIN:**

```
SELECT a.*, b.*
FROM TableA a
FULL OUTER JOIN TableB b ON a.id = b.a_id;
```

**Summary:**

- **CROSS JOIN** creates every possible pair of rows from both tables.
- **FULL OUTER JOIN** returns all rows from both tables, matching where possible, and filling with NULLs where there is no match.

---

**3. Write a SQL query to retrieve the first and last names of employees along with the names of their managers (given **Employees** and **Managers** tables).****Answer:**

To retrieve employee names along with their managers' names, you typically join the **Employees** table with the **Managers** table using a foreign key (e.g., **ManagerID** in **Employees** referencing **Managers.ManagerID**).

- **INNER JOIN:** Returns only employees who have a matching manager.
- **LEFT JOIN:** Returns all employees, including those without a manager (manager fields will be NULL).

**Example:**

```
SELECT
    e.FirstName AS EmployeeFirstName,
    e.LastName AS EmployeeLastName,
    m.FirstName AS ManagerFirstName,
    m.LastName AS ManagerLastName
FROM Employees e
LEFT JOIN Managers m ON e.ManagerID = m.ManagerID;
```

**Explanation:**

- **LEFT JOIN** is used to include employees who may not have a manager.
- **INNER JOIN** would exclude employees without a manager.

**Employees and Managers Join Result**



Join Type	Result	Use Case
INNER JOIN	Only employees with a manager are shown.	When you want to exclude employees without managers.
LEFT JOIN	All employees are shown; manager fields are NULL if no manager.	When you want to include all employees, even those without managers.

4. Write a SQL query to find the average salary for each department, given tables **Employees** (with **DepartmentID**) and **Departments** (with **DepartmentName**).

**Answer:**

To calculate the average salary for each department, join the **Employees** table with the **Departments** table on **DepartmentID**, then use **GROUP BY** to aggregate by department.

- **INNER JOIN:** Returns only departments that have at least one employee.
- **LEFT JOIN:** Returns all departments, showing **NULL** for average salary if there are no employees in a department.

**Example:**

```
SELECT
    d.DepartmentName,
    AVG(e.Salary) AS AverageSalary
FROM Departments d
LEFT JOIN Employees e ON d.DepartmentID = e.DepartmentID
GROUP BY d.DepartmentName;
```

**Explanation:**

- **LEFT JOIN** ensures all departments are listed, even those without employees.
- **AVG(e.Salary)** computes the average salary per department.
- **GROUP BY d.DepartmentName** groups results by department.

5. Write a SQL query to list all products that have never been ordered (products in a **Product** table with no matching rows in the **Orders** table).

**Answer:**

To find products that have never been ordered, you need to identify products in the **Product** table that do not have any corresponding entries in the **Orders** table. This is typically done using a **LEFT JOIN** and checking for **NULL** in the joined table, or by using a **NOT EXISTS** or **NOT IN** subquery.

- **LEFT JOIN:** Returns all products, and for those with no matching order, the order fields will be **NULL**. Filter these using **WHERE Orders.OrderID IS NULL**.
- **NOT EXISTS:** Checks for products where no matching order exists.
- **NOT IN:** Selects products whose IDs are not present in the **Orders** table.

**Example using LEFT JOIN:**

```
SELECT p.ProductID, p.ProductName
FROM Product p
LEFT JOIN Orders o ON p.ProductID = o.ProductID
WHERE o.OrderID IS NULL;
```

**Example using NOT EXISTS:**

```
SELECT p.ProductID, p.ProductName
FROM Product p
WHERE NOT EXISTS (
    SELECT 1 FROM Orders o WHERE o.ProductID = p.ProductID
);
```

**Example using NOT IN:**

```
SELECT p.ProductID, p.ProductName
FROM Product p
WHERE p.ProductID NOT IN (
    SELECT o.ProductID FROM Orders o
);
```

**Explanation:**

- **LEFT JOIN** with **WHERE o.OrderID IS NULL** finds products with no orders.
- **NOT EXISTS** and **NOT IN** are alternative approaches, often preferred for readability or performance depending on the database.
- These queries help identify products that may need promotion or removal due to lack of sales.

**Sample Data:****Sample Products Data**

ProductID	ProductName
101	Ashish's SQL Book
102	Sunil's Data Guide

If **Ashish's SQL Book** and **Sunil's Data Guide** have never been ordered, they will appear in the result.

**Approaches to Find Unordered Products**

Approach	When to Use
LEFT JOIN + IS NULL	Simple, readable, works well for moderate data sizes.
NOT EXISTS	Efficient for large datasets, especially with proper indexing.
NOT IN	Readable, but can have issues with NULLs in subquery results.

**6. Write a SQL query to list all employees who are also managers (for example, employees who appear as managers in the same table).**

**Answer:**

To find employees who are also managers, you typically use a **self-join** on the **Employees** table. This means joining the table to itself, matching employees whose **EmployeeID** appears as a **ManagerID** for other employees.

- **Self-Join:** The **Employees** table is joined to itself, using aliases to distinguish between the "employee" and the "manager" roles.
- **INNER JOIN:** Returns only those employees who are referenced as managers by at least one other employee.
- **DISTINCT:** Used to avoid duplicate rows if an employee manages multiple people.

**Example:**

```
SELECT DISTINCT
  e.EmployeeID,
  e.FirstName,
  e.LastName
FROM Employees e
INNER JOIN Employees m ON e.EmployeeID = m.ManagerID;
```

**Explanation:**

- **e** represents the employee who is also a manager.
- **m** represents employees who report to a manager.
- The join condition **e.EmployeeID = m.ManagerID** finds all employees who are listed as a manager for someone else.
- **DISTINCT** ensures each manager appears only once, even if they manage multiple employees.

**Sample Data:****Employees Table**

EmployeeID	FirstName	LastName	ManagerID
1	Ashish	Zope	NULL
2	Sunil	Patil	1
3	Ravi	Chaudhari	1
4	Ashish	Nehara	2

In this example, **Ashish Zope** (EmployeeID 1) is a manager for Sunil Patil and Ravi Chaudhari. **Sunil Patil** (EmployeeID 2) is a manager for Ashish Nehara. The query will return both Ashish Zope and Sunil Patil as employees who are also managers.

**Employees Who Are Also Managers**

EmployeeID	FirstName	LastName
1	Ashish	Zope
2	Sunil	Patil

**Summary:**

- Use a **self-join** to identify employees who are also managers.
- This pattern is common in organizational hierarchies where the manager and employee data are stored in the same table.
- The approach can be extended to retrieve additional information, such as the number of direct reports each manager has.

---

**7. What is a self-join, and when might you use it? Provide an example scenario.****Answer:**

A **self-join** is a regular join, but the table is joined with itself. This is useful when you want to compare rows within the same

table or establish relationships between rows in the same table, such as hierarchical or recursive relationships (e.g., employees and their managers).

- **Self-Join:** The same table is referenced twice in the query, using different aliases to distinguish between the two roles (e.g., employee and manager).
- **Common Use Cases:** Organizational hierarchies, bill of materials, finding pairs of related records, comparing rows within a table.

#### Example Scenario:

Suppose you have an **Employees** table where each employee may have a manager, and both employees and managers are stored in the same table.

**Employees Table**

EmployeeID	FirstName	LastName	ManagerID
1	Ashish	Zope	NULL
2	Sunil	Patil	1
3	Ravi	Chaudhari	1
4	Ashish	Nehara	2

#### Example Query:

To list each employee along with their manager's name, you can use a self-join:

```
SELECT
    e.FirstName AS EmployeeFirstName,
    e.LastName AS EmployeeLastName,
    m.FirstName AS ManagerFirstName,
    m.LastName AS ManagerLastName
FROM Employees e
LEFT JOIN Employees m ON e.ManagerID = m.EmployeeID;
```

#### Explanation:

- **e** is the alias for the employee.
- **m** is the alias for the manager.
- The join condition **e.ManagerID = m.EmployeeID** links each employee to their manager.
- **LEFT JOIN** ensures employees without a manager (e.g., Ashish Zope) are included, with manager fields as NULL.

**Employee and Manager Self-Join Result**

Employee	Manager
Ashish Zope	NULL
Sunil Patil	Ashish Zope
Ravi Chaudhari	Ashish Zope
Ashish Nehara	Sunil Patil

#### Other Example Use Cases:

- Finding all pairs of employees in the same department.
- Comparing rows for duplicates or relationships within the same table.
- Hierarchical queries, such as finding all subordinates of a manager.

### Self-Join vs Regular Join

Join Type	Purpose	Example
Self-Join	Relate rows within the same table (e.g., employee-manager relationship)	List employees and their managers using <b>Employees</b> table
Regular Join	Relate rows between different tables	Join <b>Employees</b> and <b>Departments</b> to get department names

#### Summary:

- A **self-join** is a powerful tool for querying hierarchical or related data within the same table.
- It is commonly used for organizational charts, bill of materials, and other recursive relationships.
- Use table aliases to clearly distinguish the roles of each instance of the table in the query.

---

### 8. How would you join more than two tables in a single SQL query? What factors affect the performance when joining multiple tables?

#### Answer:

Joining more than two tables in a single SQL query is common in real-world scenarios, such as retrieving employee details along with their department and manager information. This is achieved by chaining multiple **JOIN** clauses together, each connecting two tables at a time.

- **Multiple Joins:** You can join as many tables as needed by specifying additional **JOIN** clauses, using appropriate join conditions for each pair.
- **Types of Joins:** Any combination of **INNER JOIN**, **LEFT JOIN**, **RIGHT JOIN**, etc., can be used depending on the data you want to retrieve.
- **Aliases:** Table aliases help keep queries readable, especially when joining several tables.

#### Example:

Suppose you have the following tables:

- **Employees** (**EmployeeID**, **FirstName**, **LastName**, **DepartmentID**, **ManagerID**, **Salary**)
- **Departments** (**DepartmentID**, **DepartmentName**)
- **Managers** (**ManagerID**, **FirstName**, **LastName**)

To retrieve each employee's name, department, manager's name, and salary:

```
SELECT
    e.FirstName AS EmployeeFirstName,
    e.LastName AS EmployeeLastName,
    d.DepartmentName,
    m.FirstName AS ManagerFirstName,
    m.LastName AS ManagerLastName,
    e.Salary
FROM Employees e
LEFT JOIN Departments d ON e.DepartmentID = d.DepartmentID
LEFT JOIN Managers m ON e.ManagerID = m.ManagerID;
```

**Explanation:**

- **LEFT JOIN** is used to include all employees, even if they do not have a department or manager.
- Each **JOIN** connects two tables at a time, building up the result set.
- Aliases (**e**, **d**, **m**) make the query concise and readable.

**Sample Data:****Employees Table**

EmployeeID	FirstName	LastName	DepartmentID	ManagerID	Salary
1	Ashish	Zope	10	NULL	120000
2	Sunil	Patil	20	1	95000

**Departments Table**

DepartmentID	DepartmentName
10	Engineering
20	Data Science

**Managers Table**

ManagerID	FirstName	LastName
1	Ashish	Zope

**Result:****Query Result for Employee, Department, and Manager**

EmployeeFirstName	EmployeeLastName	DepartmentName	ManagerFirstName	ManagerLastName	Salary
Ashish	Zope	Engineering	NULL	NULL	120000
Sunil	Patil	Data Science	Ashish	Zope	95000

**Performance Factors When Joining Multiple Tables:****Performance Factors**

Factor	Impact	Best Practice
Indexes	Lack of indexes on join columns can cause slow queries.	Create indexes on columns used in <b>ON</b> clauses (e.g., <b>DepartmentID</b> , <b>ManagerID</b> ).
Join Order	Joining large tables first can increase intermediate result size.	Join smaller or filtered tables first when possible.
Join Type	<b>OUTER</b> joins can be slower than <b>INNER</b> joins due to more data being returned.	Use <b>INNER JOIN</b> when possible for better performance.
Data Volume	Large tables increase processing time and memory usage.	Filter data early using <b>WHERE</b> clauses.

Factor	Impact	Best Practice
Query Complexity	Complex queries with many joins can be harder to optimize.	Break down complex queries or use views for clarity.

**Summary:**

- You can join multiple tables by chaining **JOIN** clauses.
- Use table aliases for readability.
- Performance depends on indexes, join order, join type, data volume, and query complexity.
- Always test and optimize queries, especially as the number of joins increases.

### 9. Explain how an **OUTER JOIN** works when one side has no matching rows. How does this differ from an **INNER JOIN** in practice?

**Answer:**

An **OUTER JOIN** returns all rows from one (or both) tables, even if there are no matching rows in the joined table. When there is no match, the columns from the missing side are filled with **NULL** values. In contrast, an **INNER JOIN** only returns rows where there is a match in both tables.

- **LEFT OUTER JOIN (LEFT JOIN):** Returns all rows from the left table (**Employees**), and matched rows from the right table (**Departments**). If there is no match, right table columns are **NULL**.
- **RIGHT OUTER JOIN (RIGHT JOIN):** Returns all rows from the right table, and matched rows from the left table. If there is no match, left table columns are **NULL**.
- **FULL OUTER JOIN:** Returns all rows from both tables, with **NULL** in columns where there is no match.
- **INNER JOIN:** Returns only rows where there is a match in both tables.

**Example Scenario:**

Suppose you have the following tables:

**Employees Table**

EmployeeID	FirstName	LastName	DepartmentID
1	Ashish	Zope	10
2	Sunil	Patil	20
3	Ravi	Chaudhari	NULL

**Departments Table**

DepartmentID	DepartmentName
10	Engineering
20	Data Science
30	HR

**LEFT OUTER JOIN Example:**

```
SELECT
  e.FirstName,
  e.LastName,
```

```
d.DepartmentName
FROM Employees e
LEFT JOIN Departments d ON e.DepartmentID = d.DepartmentID;
```

**Result:****LEFT OUTER JOIN Result**

FirstName	LastName	DepartmentName
Ashish	Zope	Engineering
Sunil	Patil	Data Science
Ravi	Chaudhari	NULL

Notice that **Ravi Chaudhari** has no department, so **DepartmentName** is **NULL**. If you used an **INNER JOIN**, Ravi would not appear in the result.

**INNER JOIN Example:**

```
SELECT
    e.FirstName,
    e.LastName,
    d.DepartmentName
FROM Employees e
INNER JOIN Departments d ON e.DepartmentID = d.DepartmentID;
```

**Result:****INNER JOIN Result**

FirstName	LastName	DepartmentName
Ashish	Zope	Engineering
Sunil	Patil	Data Science

**Explanation:**

- **OUTER JOIN** includes all rows from one or both tables, filling in **NULL** where there is no match.
- **INNER JOIN** only includes rows where there is a match in both tables.
- Use **OUTER JOIN** when you want to see all records from one side, even if there are no matches on the other side (e.g., all employees, even those without a department).

**INNER JOIN vs OUTER JOIN**

Join Type	Rows Returned	NULLs for Missing Data?	Example Use Case
INNER JOIN	Only matching rows	No	Employees with a department
LEFT OUTER JOIN	All left table rows	Yes, for right table columns	All employees, even those without a department



Join Type	Rows Returned	NULLs for Missing Data?	Example Use Case
RIGHT OUTER JOIN	All right table rows	Yes, for left table columns	All departments, even those without employees
FULL OUTER JOIN	All rows from both tables	Yes, for missing matches on either side	All employees and all departments, showing all possible matches and non-matches

**Summary:**

- **OUTER JOIN** is useful for finding unmatched data (e.g., employees without departments, or departments without employees).
- **INNER JOIN** is used when you only care about records that exist in both tables.
- In practice, **OUTER JOIN** helps in reporting, auditing, and identifying missing relationships in your data.

## Working with Views

### 1. What is a database view, and can you update data in the base tables through it?

**Answer:**

A **view** is a virtual table based on the result of a SQL query. It does not store data itself but presents data from one or more tables. You can often update base tables through a view if the view is **simple** (e.g., based on a single table, no aggregates, no GROUP BY). Complex views (with joins, aggregates, etc.) are usually **read-only**.

- **Updatable View:** Simple SELECT from one table, no aggregates or DISTINCT.
- **Read-Only View:** Contains joins, GROUP BY, aggregate functions, or DISTINCT.

**Example:**

```
-- Simple updatable view
CREATE VIEW vw_EmployeeNames AS
SELECT EmployeeID, FirstName, LastName FROM Employees;

-- Update through the view
UPDATE vw_EmployeeNames SET FirstName = 'John' WHERE EmployeeID = 1;
```

**Summary:** Views provide a way to simplify queries and restrict access. Updates are allowed only for simple views.

### 2. What is the difference between a standard view and a materialized (or indexed) view?

#### Standard View vs Materialized View

Feature	Standard View	Materialized/Indexed View
Storage	No data stored; query runs each time	Stores result set physically
Performance	Slower for complex queries	Faster for repeated access
Refresh	Always current	Needs refresh (manual or automatic)

**Summary:** Use standard views for abstraction; use materialized views for performance on large, complex queries.

### 3. What happens if a materialized view is being refreshed (complete refresh) and a user queries it at the same time?

#### Answer:

During a **complete refresh**, the materialized view is rebuilt. If a user queries it during refresh:

- Most databases (e.g., Oracle) serve the **old data** until the refresh completes, ensuring consistency.
- Some systems may block queries or return an error if the view is unavailable.

**Tip:** Use **FAST REFRESH** or schedule refreshes during low-traffic periods to minimize impact.

---

### 4. When would you use a view in a database design? What benefits do views provide (e.g. security, abstraction)?

- **Security:** Restrict access to sensitive columns or rows.
- **Abstraction:** Hide complex joins or calculations from end users.
- **Simplification:** Provide a simple interface for reporting or applications.
- **Consistency:** Standardize business logic in one place.

#### Example:

```
-- Hide salary details from most users
CREATE VIEW vw_PublicEmployees AS
SELECT EmployeeID, FirstName, LastName FROM Employees;
```

**Summary:** Views help enforce security, simplify access, and centralize logic.

---

### 5. Can you create an index on a view? If so, what are the implications (e.g. indexed view in SQL Server)?

#### Answer:

Yes, in some databases (like SQL Server), you can create an **indexed view** (also called a materialized view). This physically stores the view's result set and creates an index on it.

- **Benefits:** Greatly improves performance for complex aggregations or joins.
- **Drawbacks:** Increases storage and maintenance overhead; restrictions on view definition (e.g., must be deterministic).

#### Example (SQL Server):

```
CREATE VIEW vw_SalesSummary WITH SCHEMABINDING AS
SELECT StoreID, SUM(SalesAmount) AS TotalSales
FROM dbo.Sales
GROUP BY StoreID;

CREATE UNIQUE CLUSTERED INDEX idx_SalesSummary_StoreID ON vw_SalesSummary(StoreID);
```

**Summary:** Indexed views boost performance but add complexity and storage cost.

---

### 6. How do you modify or drop a view if the underlying table schema changes?

- **Modify:** Use **CREATE OR REPLACE VIEW** (Oracle, PostgreSQL) or **ALTER VIEW** (SQL Server) to update the view definition.
- **Drop:** Use **DROP VIEW view\_name;** to remove the view.

- **Tip:** If a column used in the view is dropped from the base table, the view becomes invalid and must be recreated or altered.

**Example:**

```
-- Modify a view
CREATE OR REPLACE VIEW vw_EmployeeNames AS
SELECT EmployeeID, FirstName FROM Employees;

-- Drop a view
DROP VIEW vw_EmployeeNames;
```

**Summary:** Always update or drop dependent views after schema changes to base tables.

7. What is the difference between a view and a temporary table?

View vs Temporary Table

Aspect	View	Temporary Table
Persistence	Definition persists; data is always current	Exists only for session or transaction
Storage	No data stored (unless materialized)	Physically stores data
Use Case	Reusable query abstraction, security	Intermediate results, complex processing

**Summary:** Use views for abstraction and security; use temporary tables for storing intermediate results in complex queries.

Stored Procedures & Functions

1. What is the difference between a stored procedure and a user-defined function in SQL (aside from return value)?

**Answer:**

A **stored procedure** is a precompiled collection of SQL statements that can perform actions such as modifying data, controlling transactions, and returning results. A **user-defined function (UDF)** is a routine that returns a value (scalar or table) and is typically used in SELECT, WHERE, or JOIN clauses.

Stored Procedure vs User-Defined Function

Aspect	Stored Procedure	User-Defined Function
Can modify data	Yes	No (except in some DBs with special permissions)
Can be used in SELECT	No	Yes
Transaction control	Yes	No
Return type	None, scalar, or result set	Scalar or table

2. What are the advantages and disadvantages of using stored procedures?

**Answer:**

**Advantages:**

- Encapsulate business logic in the database
- Improve performance via precompilation
- Enhance security (grant EXECUTE instead of table access)
- Reduce network traffic (batch multiple statements)

**Disadvantages:**

- Harder to version and deploy than application code
- May increase database server load
- Logic split between app and DB can complicate maintenance

---

**3. Can you perform INSERT/UPDATE/DELETE operations inside a SQL function? Why or why not?****Answer:**

In most databases (e.g., SQL Server, PostgreSQL), **user-defined functions cannot perform INSERT/UPDATE/DELETE** operations on tables. This restriction ensures functions are deterministic and side-effect free, making them safe for use in queries. Some databases (like Oracle with autonomous transactions) allow limited exceptions.

---

**4. What is a table-valued function and when would you use one in a query?****Answer:**

A **table-valued function (TVF)** returns a table as its result. You use TVFs in the FROM clause of a query, similar to a regular table or view. TVFs are useful for encapsulating reusable query logic that returns sets of rows.

**Example:**

```
-- SQL Server example
CREATE FUNCTION dbo.GetEmployeesByDept(@DeptID INT)
RETURNS TABLE
AS
RETURN (
    SELECT EmployeeID, FirstName, LastName
    FROM Employees
    WHERE DepartmentID = @DeptID
);

-- Usage
SELECT * FROM dbo.GetEmployeesByDept(10);
```

---

**5. When would you use a stored procedure instead of inline SQL queries in an application?****Answer:**

Use a **stored procedure** when you want to:

- Centralize and reuse business logic
- Improve security by restricting direct table access
- Reduce SQL injection risk (parameterized execution)
- Optimize performance for complex or repetitive operations
- Batch multiple statements in a single call

---

**6. How do you pass parameters to and receive results from stored procedures?**

**Answer:**

**Parameters** are passed to stored procedures as input, output, or input/output arguments. Results can be returned via output parameters, result sets (SELECT), or return values.

**Example (SQL Server):**

```
-- Define procedure
CREATE PROCEDURE GetEmployeeByID
    @EmpID INT,
    @FirstName NVARCHAR(50) OUTPUT
AS
BEGIN
    SELECT @FirstName = FirstName FROM Employees WHERE EmployeeID = @EmpID;
END;

-- Execute procedure
DECLARE @Name NVARCHAR(50);
EXEC GetEmployeeByID 1, @Name OUTPUT;
SELECT @Name;
```

---

**7. How would you debug or test a slow or failing stored procedure in production?****Answer:**

- Capture execution plans to identify bottlenecks
- Use logging or print statements for tracing
- Test with sample data in a development environment
- Check for blocking, deadlocks, or resource waits
- Review recent schema or data changes
- Use database profiling tools (e.g., SQL Profiler, Extended Events)

---

**8. How do you grant a user permission to execute a specific stored procedure?****Answer:**

Use the **GRANT EXECUTE** statement to allow a user to run a stored procedure.

**Example:**

```
GRANT EXECUTE ON OBJECT::GetEmployeeByID TO username;
```

*This grants the user permission to execute the **GetEmployeeByID** procedure.*

---

## Triggers & Automation

---

**1. What is a trigger in SQL, and when would you use one? Give an example use case.****Answer:**

A **trigger** is a special kind of stored procedure that automatically executes in response to certain events on a table or view (such as **INSERT**, **UPDATE**, or **DELETE**). Triggers are used for enforcing business rules, auditing changes, maintaining derived data, or automating tasks.

- **Use Cases:** Auditing changes, enforcing complex constraints, cascading updates/deletes, logging, or synchronizing tables.

**Example:**

```
-- Audit trigger: Log every employee salary change
CREATE TRIGGER trg_AuditSalaryChange
ON Employees
AFTER UPDATE
AS
BEGIN
    INSERT INTO SalaryAudit (EmployeeID, OldSalary, NewSalary, ChangedAt)
    SELECT i.EmployeeID, d.Salary, i.Salary, GETDATE()
    FROM inserted i
    JOIN deleted d ON i.EmployeeID = d.EmployeeID
    WHERE i.Salary <> d.Salary;
END;
```

*This trigger logs salary changes to an audit table whenever an employee's salary is updated.*

---

**2. What is the difference between an AFTER trigger and an INSTEAD OF trigger (e.g. in SQL Server)?****3. What are the “inserted” and “deleted” magic tables in SQL Server triggers?****Answer:**

In SQL Server, **inserted** and **deleted** are special (magic) tables available inside triggers:

- **inserted:** Holds the new rows for **INSERT** and **UPDATE** operations.
- **deleted:** Holds the old rows for **DELETE** and **UPDATE** operations.

**Example:** In an **AFTER UPDATE** trigger, **inserted** has new values, **deleted** has old values.

---

**4. How can triggers be used to enforce business rules or data integrity (e.g. auditing changes, simulating foreign keys)?****Answer:**

Triggers can enforce business rules by validating data, preventing invalid changes, logging modifications, or simulating constraints not natively supported (e.g., cascading deletes, custom referential integrity).

- **Auditing:** Log changes to sensitive data (e.g., salary updates).
- **Enforcing Rules:** Prevent deletion of parent rows if child rows exist (simulate foreign key).
- **Data Integrity:** Automatically update related tables or maintain derived columns.

**Example:**

```
-- Prevent deleting a department if employees exist
CREATE TRIGGER trg_PreventDeptDelete
ON Departments
INSTEAD OF DELETE
AS
BEGIN
    IF EXISTS (SELECT 1 FROM Employees e JOIN deleted d ON e.DepartmentID = d.DepartmentID)
        RAISERROR('Cannot delete department with employees.', 16, 1);
    ELSE
```

```
DELETE FROM Departments WHERE DepartmentID IN (SELECT DepartmentID FROM deleted);
END;
```

*This trigger blocks deletion of a department if employees are assigned to it.*

## 5. What are the potential drawbacks of using triggers (such as performance impact or hidden logic)?

### Answer:

Triggers can introduce hidden logic and performance overhead:

- **Performance:** Triggers add extra processing to DML operations, potentially slowing down inserts, updates, or deletes.
- **Hidden Logic:** Business rules in triggers may not be obvious to developers, making debugging and maintenance harder.
- **Complexity:** Nested or recursive triggers can cause unexpected behavior.
- **Portability:** Trigger syntax and behavior can vary between database systems.

**Summary:** Use triggers judiciously; document their behavior and monitor performance.

## 6. How do INSTEAD OF triggers on a view work?

### Answer:

**INSTEAD OF triggers** on a view intercept **INSERT**, **UPDATE**, or **DELETE** operations and allow you to define custom logic for how those operations are handled. This is useful for updatable views that join multiple tables or require special handling.

### Example:

```
-- Allow updates to a view that joins Employees and Departments
CREATE VIEW vw_EmployeeDept AS
SELECT e.EmployeeID, e.FirstName, d.DepartmentName
FROM Employees e
JOIN Departments d ON e.DepartmentID = d.DepartmentID;

CREATE TRIGGER trg_UpdateEmpDept
ON vw_EmployeeDept
INSTEAD OF UPDATE
AS
BEGIN
    UPDATE Employees
    SET FirstName = i.FirstName
    FROM inserted i
    WHERE Employees.EmployeeID = i.EmployeeID;
END;
```

*This trigger allows updates to the **FirstName** column via the view.*

## 7. Can triggers call stored procedures, and are there any limitations to doing that?

### Answer:

Yes, triggers can call stored procedures. However, there are limitations:

- **Side Effects:** Procedures called from triggers should not commit/rollback transactions independently.
- **Performance:** Long-running procedures can slow down DML operations.

- **Recursion:** Be careful to avoid recursive trigger/procedure calls.
- **Permissions:** The trigger must have permission to execute the procedure.

**Summary:** Triggers can call stored procedures, but keep logic efficient and avoid transactional conflicts.

---

## Transactions & Concurrency Control

---

### 1. What are the ACID properties of a database transaction (atomicity, consistency, isolation, durability)?

**Answer:**

**ACID** stands for:

- **Atomicity:** All operations in a transaction succeed or all fail (no partial changes).
- **Consistency:** Transactions bring the database from one valid state to another, preserving rules.
- **Isolation:** Concurrent transactions do not interfere; intermediate states are hidden.
- **Durability:** Once committed, changes are permanent even after a crash.

---

### 2. What are the different SQL isolation levels (READ UNCOMMITTED, READ COMMITTED, REPEATABLE READ, SERIALIZABLE), and what phenomena do they prevent (dirty reads, non-repeatable reads, phantom reads)?

SQL Isolation Levels and Phenomena

Level	Dirty Reads	Non-Repeatable Reads	Phantom Reads
READ UNCOMMITTED	Possible	Possible	Possible
READ COMMITTED	Prevented	Possible	Possible
REPEATABLE READ	Prevented	Prevented	Possible
SERIALIZABLE	Prevented	Prevented	Prevented

**Summary:** Higher isolation = fewer anomalies, but more locking and lower concurrency.

---

### 3. What is a deadlock in database terms, and how can you prevent or resolve deadlocks?

**Answer:**

A **deadlock** occurs when two or more transactions block each other, each waiting for the other to release locks. Neither can proceed.

- **Prevention:** Access tables in the same order, keep transactions short, use lower isolation levels if possible.
- **Resolution:** The database detects deadlocks and aborts (rolls back) one transaction (the "victim").

---

### 4. How do you control transactions in SQL (BEGIN, COMMIT, ROLLBACK)? Give an example of using a transaction in a stored procedure or batch.

**Answer:**

Use **BEGIN TRANSACTION** to start, **COMMIT** to save, and **ROLLBACK** to undo.

**Example:**



```

BEGIN TRANSACTION;
    UPDATE Accounts SET Balance = Balance - 100 WHERE AccountID = 1;
    UPDATE Accounts SET Balance = Balance + 100 WHERE AccountID = 2;
IF @@ERROR <> 0
    ROLLBACK;
ELSE
    COMMIT;

```

*This ensures both updates succeed or both are undone.*

---

#### 5. If you run a long SELECT query on a table while another transaction is updating rows in that table, will your session see the old data or new data by default? (Consider default isolation level behavior.)

##### Answer:

By default (READ COMMITTED), your SELECT sees only committed data at the time each row is read. You may see new data if the update commits before your SELECT reads that row.

---

#### 6. What is the difference between pessimistic and optimistic locking, and when would you use each?

- **Pessimistic Locking:** Locks data when read, blocking others until transaction ends. Use for high-conflict scenarios.
- **Optimistic Locking:** No locks when reading; checks for changes before writing (e.g., using a version column). Use when conflicts are rare.

---

#### 7. What is a savepoint in a transaction, and how do you use it?

##### Answer:

A **savepoint** marks a point within a transaction to which you can roll back without affecting earlier work.

##### Example:

```

BEGIN TRANSACTION;
    UPDATE Accounts SET Balance = Balance - 100 WHERE AccountID = 1;
    SAVE TRANSACTION Save1;
    UPDATE Accounts SET Balance = Balance + 100 WHERE AccountID = 2;
    -- If needed:
    ROLLBACK TRANSACTION Save1;
COMMIT;

```

---

#### 8. How do two-phase commit protocols work in distributed transactions?

##### Answer:

**Two-phase commit** ensures all participants in a distributed transaction agree to commit or roll back:

1. **Prepare phase:** Coordinator asks all nodes if they can commit.
2. **Commit phase:** If all agree, coordinator tells all to commit; otherwise, tells all to roll back.

---

#### 9. How can you identify and terminate a blocking or long-running transaction in a SQL database?

- **Identify:** Use system views (e.g., `sys.dm_exec_requests`, `sp_who2` in SQL Server) to find blocking sessions.
- **Terminate:** Use `KILL session_id` (SQL Server) or `ALTER SYSTEM KILL SESSION` (Oracle) to end the session.

10. What is deadlock detection, and how does the database engine choose a deadlock victim?

Answer:

The database periodically checks for deadlocks. When found, it picks a "victim" transaction to roll back (usually the one with the least cost or least work done) to break the cycle and let others proceed.

## Performance Tuning & Query Optimization

### Performance Tuning & Query Optimization: Interview Questions & Answers

1. What is a query execution plan and how do you use it to improve performance?

Answer:

A **query execution plan** is a detailed roadmap generated by the database engine that outlines how a SQL query will be executed. It includes the sequence of operations (such as scans, joins, sorts), the access methods (e.g., index seek, table scan), and the estimated cost of each step. By analyzing the execution plan, you can identify inefficiencies like full table scans, missing indexes, or suboptimal join strategies, and then optimize your queries or schema accordingly.

- **How to View:** Use `EXPLAIN` (MySQL, PostgreSQL), `EXPLAIN PLAN` (Oracle), or graphical tools (SQL Server Management Studio, pgAdmin, etc.).
- **Optimization:** Look for expensive operations (e.g., table scans), add or adjust indexes, rewrite queries, or refactor schema based on plan insights.

#### Common Execution Plan Operators

Operator	Description
Table Scan	Reads all rows in a table; slow for large tables
Index Seek	Efficiently finds rows using an index
Nested Loop Join	Efficient for small result sets or indexed joins
Hash Join	Good for joining large, unsorted data sets

Syntax & Example:

```
EXPLAIN SELECT * FROM Employees WHERE DepartmentID = 10;
```

This command shows the execution plan for the query, including whether an index is used.

**Summary:** Always review execution plans for slow queries to identify and resolve performance bottlenecks.

2. How would you optimize a slow SQL query in production?

Answer:

Optimizing a slow query involves a systematic approach:

- **Analyze the Execution Plan:** Identify bottlenecks such as full table scans, missing indexes, or expensive joins.
- **Add Indexes:** Create indexes on columns used in WHERE, JOIN, and ORDER BY clauses.
- **Rewrite Queries:** Simplify complex queries, avoid unnecessary subqueries, and use set-based operations instead of row-by-row processing.
- **Select Only Needed Columns:** Avoid `SELECT *`; retrieve only the columns you need.

- **Partition Large Tables:** Use table partitioning to improve query performance on very large datasets.
- **Update Statistics & Rebuild Indexes:** Ensure the query optimizer has up-to-date statistics and that indexes are not fragmented.

Optimization Techniques Comparison

Technique	When to Use
Indexing	Frequent filtering or joining on columns
Query Rewrite	Complex or inefficient queries
Partitioning	Very large tables

**Summary:** Use a combination of indexing, query rewriting, and regular maintenance for optimal performance.

3. What is the difference between UNION and UNION ALL in SQL, and when would you use each?

**Answer:**

**UNION** combines the results of two queries and removes duplicate rows. **UNION ALL** combines results and keeps all duplicates. **UNION ALL** is faster because it does not perform the extra step of removing duplicates.

UNION vs UNION ALL

Operator	Duplicates Removed?	Performance
UNION	Yes	Slower (deduplication required)
UNION ALL	No	Faster

**Syntax & Example:**

```
SELECT Name FROM Employees
UNION
SELECT Name FROM Managers;

SELECT Name FROM Employees
UNION ALL
SELECT Name FROM Managers;
```

**Summary:** Use **UNION ALL** for better performance if you do not need to remove duplicates.

4. How can you find duplicate rows in a table using SQL?

**Answer:**

To identify duplicate rows in a table, you typically use the **GROUP BY** clause on the columns that define a duplicate, combined with the **HAVING** clause to filter groups that occur more than once. This approach helps you find which values are repeated and how many times they appear.

- **Step 1:** Decide which columns define a duplicate (e.g., all columns, or a subset such as `email` or `first_name, last_name`).
- **Step 2:** Use **GROUP BY** on those columns and count the occurrences.
- **Step 3:** Use **HAVING COUNT(\*) > 1** to filter only the duplicates.

**Syntax & Example:**

```
-- Example: Find duplicates based on column1 and column2
SELECT column1, column2, COUNT(*) AS duplicate_count
FROM table_name
GROUP BY column1, column2
HAVING COUNT(*) > 1;
```

This query lists each combination of *column1* and *column2* that appears more than once, along with the number of times it occurs.

**Example Scenario:**

Suppose you have a *users* table with columns *email* and *username*, and you want to find duplicate emails:

```
SELECT email, COUNT(*) AS duplicate_count
FROM users
GROUP BY email
HAVING COUNT(*) > 1;
```

This returns all email addresses that appear more than once in the *users* table.

**Sample Output**

email	duplicate_count
ashish@example.com	3
sunil@example.com	2

**To retrieve the full duplicate rows** (not just the duplicate values), you can join the results back to the original table:

```
SELECT t.*
FROM table_name t
JOIN (
    SELECT column1, column2
    FROM table_name
    GROUP BY column1, column2
    HAVING COUNT(*) > 1
) dup
ON t.column1 = dup.column1 AND t.column2 = dup.column2;
```

This returns all rows from *table\_name* that are considered duplicates based on *column1* and *column2*.

**Summary:** Use **GROUP BY** and **HAVING COUNT(\*) > 1** to find duplicates. Adjust the columns in **GROUP BY** to match your definition of a duplicate row.

---

**5. Write a SQL query to find the 10th highest salary in an Employee table.****Answer:**

To find the 10th highest salary, you need to rank the unique salary values in descending order and select the one at

position 10. There are several approaches, depending on your SQL dialect and requirements (e.g., handling duplicate salaries).

- **Approach 1: Subquery with DISTINCT, ORDER BY, LIMIT/OFFSET (MySQL, PostgreSQL)**

This method selects the top 10 unique salaries in descending order, then picks the minimum (which is the 10th highest).

```
SELECT MIN(Salary) AS TenthHighestSalary
FROM (
    SELECT DISTINCT Salary
    FROM Employee
    ORDER BY Salary DESC
    LIMIT 10
) AS Top10;
```

*This returns the 10th highest unique salary. If there are fewer than 10 unique salaries, it returns NULL.*

- **Approach 2: Using OFFSET (MySQL, PostgreSQL)**

You can also use **OFFSET** to directly get the 10th highest salary:

```
SELECT DISTINCT Salary
FROM Employee
ORDER BY Salary DESC
LIMIT 1 OFFSET 9;
```

*This returns the 10th highest unique salary (OFFSET is zero-based).*

- **Approach 3: Using Window Functions (SQL Server, PostgreSQL, Oracle)**

Window functions like **DENSE\_RANK()** can be used to assign a rank to each unique salary:

```
SELECT Salary AS TenthHighestSalary
FROM (
    SELECT Salary, DENSE_RANK() OVER (ORDER BY Salary DESC) AS rnk
    FROM Employee
) t
WHERE rnk = 10;
```

*This works in databases that support window functions and handles ties (duplicate salaries) correctly.*

- **Approach 4: Correlated Subquery (Standard SQL)**

This approach counts how many distinct salaries are greater than the current one:

```
SELECT Salary AS TenthHighestSalary
FROM Employee e1
WHERE (
    SELECT COUNT(DISTINCT Salary)
    FROM Employee e2
    WHERE e2.Salary > e1.Salary
) = 9
LIMIT 1;
```

*This returns the 10th highest unique salary. If there are fewer than 10 unique salaries, it returns no row.*

**Summary:** Use `DISTINCT` to ignore duplicate salaries. Choose the approach that fits your SQL dialect and performance needs. Window functions are preferred for clarity and efficiency in modern databases.

---

## 6. How would you retrieve the last 5 records (by date or ID) from a table?

### Answer:

To retrieve the last 5 records from a table, you need to determine which column defines the "order" of your data—typically a date column (e.g., `created_at`, `order_date`) or a unique, incrementing ID (e.g., `id`, `record_id`). You then sort the table in descending order by that column and use the `LIMIT` clause (or its equivalent) to fetch only the top 5 rows. This approach works in most SQL databases, including MySQL, PostgreSQL, and SQLite.

- **Step 1:** Identify the column that determines the record order (e.g., `date_column` or `id`).
- **Step 2:** Use `ORDER BY` in descending order (`DESC`) to bring the latest or highest values to the top.
- **Step 3:** Use `LIMIT 5` to restrict the result to the last 5 records.

### Syntax & Examples:

- **By Date:**

```
SELECT *
FROM table_name
ORDER BY date_column DESC
LIMIT 5;
```

*This returns the 5 most recent records based on `date_column`.*

- **By ID (assuming higher IDs are newer):**

```
SELECT *
FROM table_name
ORDER BY id DESC
LIMIT 5;
```

*This returns the 5 records with the highest (latest) IDs.*

- **To return the last 5 records in ascending order (oldest to newest among the last 5):**

```
SELECT *
FROM (
    SELECT *
    FROM table_name
    ORDER BY date_column DESC
    LIMIT 5
) AS last_five
ORDER BY date_column ASC;
```

*This subquery first selects the last 5 records, then reorders them in ascending order.*

Database	Syntax
MySQL / PostgreSQL / SQLite	<code>ORDER BY column DESC LIMIT 5</code>
SQL Server	<code>SELECT TOP 5 * FROM table_name ORDER BY column DESC</code>
Oracle	<code>SELECT * FROM (SELECT * FROM table_name ORDER BY column DESC) WHERE ROWNUM &lt;= 5</code>

**Summary:** Use `ORDER BY ... DESC LIMIT 5` to get the last 5 records by date or ID. Adjust the column and syntax for your database system.

---

## 7. Write a SQL query to exclude specific values (e.g., select all rows except those where ID is X or Y).

### Answer:

To exclude specific values from your query results, use the `NOT IN` operator or multiple `!=` (or `<>`) conditions in the `WHERE` clause. This is useful when you want to filter out rows with certain values in a column, such as excluding students with specific IDs.

- **NOT IN:** Excludes all rows where the column matches any value in the list.
- **!= or <>:** Excludes rows matching a single value; combine with `AND` for multiple exclusions.
- **NOT EQUALS (for a single value):** Use `WHERE ID != X` or `WHERE ID <> X`.

### Syntax & Examples:

- **Exclude multiple values using NOT IN:**

```
SELECT *
FROM Student
WHERE ID NOT IN (101, 102);
```

*This query returns all students except those with ID 101 or 102.*

- **Exclude a single value using != (or <>):**

```
SELECT *
FROM Student
WHERE ID != 101;
-- or
SELECT *
FROM Student
WHERE ID <> 101;
```

*This returns all students except the one with ID 101.*

- **Exclude multiple values using AND:**

```
SELECT *
FROM Student
WHERE ID != 101 AND ID != 102;
```

This is equivalent to using `NOT IN (101, 102)`.

- **Exclude values in a string column:**

```
SELECT *
FROM Student
WHERE Name NOT IN ('Ashish', 'Sunil');
```

This excludes students named Ashish or Sunil.

#### Exclusion Methods Comparison

Method	Use Case	Example
NOT IN	Exclude multiple values	<code>ID NOT IN (101, 102)</code>
<code>!=</code> or <code>&lt;&gt;</code>	Exclude a single value	<code>ID != 101</code>
AND with <code>!=</code>	Exclude several values (few)	<code>ID != 101 AND ID != 102</code>

#### Notes:

- If the column contains `NULL` values, `NOT IN` may return no rows if any value in the list is `NULL`. Use `IS NOT NULL` if needed.
- For large exclusion lists, `NOT IN` is more concise and readable.

**Summary:** Use `NOT IN` to exclude multiple values, or `!=`/`<>` for single values. Adjust the column and values as needed for your query.

---

### 8. How do you retrieve the Nth record (e.g., the 3rd record) from a table?

#### Answer:

Retrieving the Nth record from a table depends on the SQL dialect and whether you want the Nth row in a specific order. The most common approach is to use `ORDER BY` with `LIMIT` and `OFFSET` (supported in MySQL, PostgreSQL, SQLite), or `ROW_NUMBER()` window function (in SQL Server, PostgreSQL, Oracle).

- **Step 1:** Decide the column(s) that define the order (e.g., `created_at`, `id`).
- **Step 2:** Use `ORDER BY` to specify the order.
- **Step 3:** Use `OFFSET` to skip the first N-1 rows and `LIMIT 1` to get the Nth row.

#### Syntax & Examples:

- **MySQL / PostgreSQL / SQLite:**

```
SELECT *
FROM table_name
ORDER BY ordering_column
LIMIT 1 OFFSET N-1;
```

For example, to get the 3rd record (N=3):

```
SELECT *
FROM table_name
```



```
ORDER BY ordering_column
LIMIT 1 OFFSET 2;
```

This returns the 3rd row in the specified order (OFFSET is zero-based).

- **SQL Server:**

```
SELECT *
FROM (
    SELECT *, ROW_NUMBER() OVER (ORDER BY ordering_column) AS rn
    FROM table_name
) t
WHERE rn = N;
```

Replace **N** with the desired row number (e.g., 3 for the 3rd record).

- **Oracle:**

```
SELECT *
FROM (
    SELECT t.*, ROW_NUMBER() OVER (ORDER BY ordering_column) AS rn
    FROM table_name t
)
WHERE rn = N;
```

### Nth Record Retrieval Methods by SQL Dialect

Database	Syntax
MySQL / PostgreSQL / SQLite	ORDER BY ... LIMIT 1 OFFSET N-1
SQL Server / Oracle	ROW_NUMBER() OVER (ORDER BY ...) = N

#### Notes:

- If the table has no explicit ordering, the result may be unpredictable. Always use **ORDER BY** for deterministic results.
- For large tables, ensure the ordering column is indexed for better performance.

**Summary:** Use **ORDER BY** with **LIMIT 1 OFFSET N-1** (or **ROW\_NUMBER()** in SQL Server/Oracle) to retrieve the Nth record in a specified order.

## 9. How do you obtain the CREATE TABLE DDL for an existing table in SQL?

### Answer:

To obtain the **CREATE TABLE** DDL (Data Definition Language) statement for an existing table, you use database-specific commands or tools that generate the SQL statement required to recreate the table structure, including columns, data types, constraints, indexes, and other properties. This is useful for documentation, migration, backup, or recreating tables in another environment.

### DDL Retrieval Methods by Database

Database	How to Get CREATE TABLE DDL	Example
----------	-----------------------------	---------

Database	How to Get CREATE TABLE DDL	Example
MySQL	Use <code>SHOW CREATE TABLE</code> command	<code>SHOW CREATE TABLE employees;</code>
PostgreSQL	Use <code>pg_dump</code> with <code>--schema-only</code> or query <code>pg_catalog/information_schema</code> . Many GUI tools (e.g., pgAdmin) also provide "Generate SQL" or "DDL" options.	<code>pg_dump -U username -d dbname -t employees --schema-only</code>
SQL Server	Use SQL Server Management Studio (SSMS): right-click table → Script Table as → CREATE To → New Query Editor Window. Or use <code>sp_help</code> for table details.	<i>SSMS GUI: Script Table as → CREATE To</i>
Oracle	Use <code>DBMS_METADATA.GET_DDL</code> function or tools like SQL Developer.	<code>SELECT DBMS_METADATA.GET_DDL('TABLE', 'EMPLOYEES') FROM DUAL;</code>
SQLite	Query <code>sqlite_master</code> table for the SQL statement.	<code>SELECT sql FROM sqlite_master WHERE type='table' AND name='employees';</code>

- **GUI Tools:** Most database management tools (e.g., MySQL Workbench, pgAdmin, SSMS, Oracle SQL Developer) provide right-click options to generate the CREATE TABLE script for any table.
- **Command-Line Utilities:** Utilities like `mysqldump`, `pg_dump`, or `expdp` (Oracle) can export DDL for tables or entire schemas.
- **Information Schema:** For advanced scripting, you can query metadata tables (e.g., `information_schema.columns`) to reconstruct DDL, but this is rarely needed for standard use cases.

#### Example Output (MySQL):

```
SHOW CREATE TABLE employees;
```

#### SHOW CREATE TABLE employees Result

Table	Create Table
employees	<pre>CREATE TABLE `employees` (   `id` int NOT NULL AUTO_INCREMENT,   `first_name` varchar(50) NOT NULL,   `last_name` varchar(50) NOT NULL,   PRIMARY KEY (`id`) ) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4;</pre>

**Summary:** Use the appropriate command or tool for your database system to generate the CREATE TABLE DDL for an existing table. This is essential for migrations, backups, and documentation.

---

## 10. Explain the difference between the RANK() and DENSE\_RANK() window functions.

### Answer:

**RANK()** and **DENSE\_RANK()** are window functions used to assign a ranking number to each row within a result set, based

on the ordering of one or more columns. They are commonly used for tasks like leaderboard generation, top-N queries, and reporting. The key difference between them is how they handle ties (rows with equal values in the ordering column).

- **RANK()**: Assigns the same rank to tied rows, but leaves gaps in the ranking sequence after the tie. The next rank after a tie is incremented by the number of tied rows.
- **DENSE\_RANK()**: Also assigns the same rank to tied rows, but does not leave gaps. The next rank after a tie is incremented by one, regardless of the number of tied rows.

**Syntax:**

```
SELECT
    column1,
    RANK() OVER (ORDER BY column2 DESC) AS rank,
    DENSE_RANK() OVER (ORDER BY column2 DESC) AS dense_rank
FROM table_name;
```

**Example:**

Suppose you have a **Scores** table:

Player	Score
A	100
B	90
C	90
D	80

Applying **RANK()** and **DENSE\_RANK()**:

Player	Score	RANK()	DENSE_RANK()
A	100	1	1
B	90	2	2
C	90	2	2
D	80	4	3

- With **RANK()**, both B and C have rank 2 (tie), and the next rank is 4 (skips 3).
- With **DENSE\_RANK()**, both B and C have rank 2, and the next rank is 3 (no gap).

**RANK() vs DENSE\_RANK() Comparison**

Function	Ranking Behavior	Example Output
RANK()	Skips ranks after ties (gaps in ranking)	1, 2, 2, 4
DENSE_RANK()	No gaps after ties (consecutive ranking)	1, 2, 2, 3

**When to use:**

- Use **RANK()** when you want to reflect the number of tied rows in the ranking (e.g., for competition scoring where the next rank skips ahead).

- Use **DENSE\_RANK()** when you want consecutive ranking numbers without gaps, even if there are ties (e.g., for reporting or grouping).

**Summary:** **RANK()** leaves gaps after ties; **DENSE\_RANK()** does not. Choose based on your ranking requirements.

---

## 11. When would you use **ROW\_NUMBER()**, **RANK()**, or **DENSE\_RANK()** in a query? Give a use case.

### Answer:

**ROW\_NUMBER()**, **RANK()**, and **DENSE\_RANK()** are window functions used to assign a sequential number or rank to rows within a result set, based on a specified ordering. They are commonly used for tasks such as pagination, ranking, deduplication, and leaderboard generation. The choice among them depends on how you want to handle ties (rows with equal values in the ordering column).

- **ROW\_NUMBER():** Assigns a unique sequential number to each row within the partition, regardless of ties. No two rows get the same number, even if their values are identical.
  - **Use Case:** Pagination (fetching rows N to M), removing duplicates (keeping only the first occurrence), or selecting the "top N" per group.
- **RANK():** Assigns the same rank to tied rows, but leaves gaps in the ranking sequence after ties. The next rank after a tie is incremented by the number of tied rows.
  - **Use Case:** Competition ranking where ties should skip ranks (e.g., Olympic medals: 1, 2, 2, 4).
- **DENSE\_RANK():** Assigns the same rank to tied rows, but does not leave gaps. The next rank after a tie is incremented by one.
  - **Use Case:** Leaderboards or reporting where consecutive ranking is desired without gaps (e.g., 1, 2, 2, 3).

### Syntax & Example:

Suppose you have a **Sales** table:

Employee	SalesAmount
Alice	500
Bob	400
Carol	400
Dave	300

To assign rankings based on **SalesAmount** (highest first):

```
SELECT
    Employee,
    SalesAmount,
    ROW_NUMBER() OVER (ORDER BY SalesAmount DESC) AS row_num,
    RANK() OVER (ORDER BY SalesAmount DESC) AS rank,
    DENSE_RANK() OVER (ORDER BY SalesAmount DESC) AS dense_rank
FROM Sales;
```

Result:

Employee	SalesAmount	row_num	rank	dense_rank
Alice	500	1	1	1
Bob	400	2	2	2

Employee	SalesAmount	row_num	rank	dense_rank
Carol	400	3	2	2
Dave	300	4	4	3

- **ROW\_NUMBER():** Each row gets a unique number (no ties).
- **RANK():** Bob and Carol tie for 2nd place, so both get rank 2; the next rank is 4 (gap).
- **DENSE\_RANK():** Bob and Carol tie for 2nd, both get 2; the next rank is 3 (no gap).

#### Window Function Use Cases

Function	Typical Use Case
ROW_NUMBER()	Pagination, deduplication, selecting Nth row
RANK()	Competition ranking with gaps after ties
DENSE_RANK()	Leaderboard or reporting with consecutive ranks

**Summary:** Use `ROW_NUMBER()` for unique row numbering, `RANK()` for rankings with gaps after ties, and `DENSE_RANK()` for consecutive rankings without gaps. Choose based on your business logic and how you want to handle ties.

## 12. Write a SQL query to compute the median number of searches made by users, given a summary table of search counts.

### Answer:

The **median** is the middle value in a sorted list of numbers. If the count of values is odd, the median is the value at the center. If the count is even, the median is the average of the two central values. Calculating the median in SQL requires ordering the data and identifying the middle value(s). This is typically done using window functions such as `ROW_NUMBER()`, `RANK()`, or `PERCENTILE_CONT()` (if supported by your database).

- **Step 1:** Order the search counts for all users.
- **Step 2:** Assign a row number to each record and count the total number of records.
- **Step 3:** Select the middle row(s) based on the total count.
- **Step 4:** For even counts, average the two middle values; for odd counts, select the single middle value.

### Syntax & Example (Standard SQL):

Suppose you have a table `user_search_summary` with columns `user_id` and `search_count`.

```
SELECT AVG(search_count) AS median_searches
FROM (
  SELECT
    search_count,
    ROW_NUMBER() OVER (ORDER BY search_count) AS rn,
    COUNT(*) OVER () AS total
  FROM user_search_summary
) t
WHERE
  rn = (total + 1) / 2
  OR (total % 2 = 0 AND rn = (total / 2) + 1);
```

- This query works for both odd and even numbers of rows.
- For odd counts, it selects the middle row. For even counts, it averages the two central rows.

- `ROW_NUMBER()` assigns a unique sequential number to each row in order of `search_count`.
- `COUNT(*) OVER ()` gives the total number of rows.

**Alternative (Using PERCENTILE\_CONT, if supported):**

Some databases (e.g., PostgreSQL, Oracle, SQL Server 2012+) support the `PERCENTILE_CONT` window function, which directly computes the median:

```
SELECT
    PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY search_count) AS median_searches
FROM user_search_summary;
```

- This is the most concise and efficient way to compute the median if your database supports it.

**Sample Data:**

user_id	search_count
1	5
2	2
3	7
4	3
5	4

- Ordered `search_count`: 2, 3, 4, 5, 7
- Median is 4 (the third value in the sorted list)

**Summary:** To compute the median in SQL, use window functions to rank and select the middle value(s), or use `PERCENTILE_CONT(0.5)` if available. The median provides a robust measure of central tendency, especially when data contains outliers.

**13. Write a SQL query to calculate the sum of odd-numbered and even-numbered measurements separately for each day.****Answer:**

To calculate the sum of odd-numbered and even-numbered measurements for each day, you can use a **CASE** expression inside the **SUM()** aggregate function. This allows you to conditionally sum values based on whether the `measurement_number` is odd or even. The `MOD()` (or `%` in some databases) function is used to determine if a number is odd (`MOD(measurement_number, 2) = 1`) or even (`MOD(measurement_number, 2) = 0`).

- **Step 1:** Use `CASE` to check if `measurement_number` is odd or even.
- **Step 2:** Use `SUM()` to aggregate the values for odd and even numbers separately.
- **Step 3:** `GROUP BY` the `day` column to get results per day.

**Syntax & Example (Standard SQL):**

```
SELECT
    day,
    SUM(CASE WHEN MOD(measurement_number, 2) = 1 THEN value ELSE 0 END) AS odd_sum,
    SUM(CASE WHEN MOD(measurement_number, 2) = 0 THEN value ELSE 0 END) AS even_sum
FROM measurements
GROUP BY day;
```

- `MOD(measurement_number, 2) = 1` checks for odd numbers.
- `MOD(measurement_number, 2) = 0` checks for even numbers.
- Replace `MOD` with `%` if your database uses that syntax (e.g., `measurement_number % 2`).

**Example Data:**

day	measurement_number	value
2024-06-01	1	10
2024-06-01	2	20
2024-06-01	3	30
2024-06-02	1	15
2024-06-02	2	25

**Result:**

day	odd_sum	even_sum
2024-06-01	40	20
2024-06-02	15	25

- For 2024-06-01: odd measurements (1 and 3) sum to 10 + 30 = 40; even measurement (2) is 20.
- For 2024-06-02: odd measurement (1) is 15; even measurement (2) is 25.

**Summary:** Use `CASE` with `SUM()` and `MOD()` (or `%`) to separate and aggregate odd and even measurement values for each day.

**14. Write a SQL query to get the average review rating for each product for each month.**

**Answer:**

To calculate the average review rating for each product for each month, you need to:

- Group reviews by **product** and by **month** (derived from the review date).
- Use an aggregate function (`AVG()`) to compute the average rating per group.
- Use a date truncation or formatting function to extract the month from the review date. The exact function depends on your SQL dialect.

**Date Truncation Functions by SQL Dialect**

Database	Function	Example
PostgreSQL	<code>DATE_TRUNC('month', review_date)</code>	<code>DATE_TRUNC('month', review_date)</code>
MySQL	<code>DATE_FORMAT(review_date, '%Y-%m')</code>	<code>DATE_FORMAT(review_date, '%Y-%m')</code>
SQL Server	<code>FORMAT(review_date, 'yyyy-MM')</code> or <code>YEAR(review_date), MONTH(review_date)</code>	<code>FORMAT(review_date, 'yyyy-MM')</code>
Oracle	<code>TO_CHAR(review_date, 'YYYY-MM')</code>	<code>TO_CHAR(review_date, 'YYYY-MM')</code>

**Syntax & Examples:****• PostgreSQL:**

```
SELECT
    product_id,
    DATE_TRUNC('month', review_date) AS review_month,
    AVG(rating) AS avg_rating
FROM reviews
GROUP BY product_id, DATE_TRUNC('month', review_date)
ORDER BY product_id, review_month;
```

**• MySQL:**

```
SELECT
    product_id,
    DATE_FORMAT(review_date, '%Y-%m') AS review_month,
    AVG(rating) AS avg_rating
FROM reviews
GROUP BY product_id, review_month
ORDER BY product_id, review_month;
```

**• SQL Server:**

```
SELECT
    product_id,
    FORMAT(review_date, 'yyyy-MM') AS review_month,
    AVG(rating) AS avg_rating
FROM reviews
GROUP BY product_id, FORMAT(review_date, 'yyyy-MM')
ORDER BY product_id, review_month;
```

**• Oracle:**

```
SELECT
    product_id,
    TO_CHAR(review_date, 'YYYY-MM') AS review_month,
    AVG(rating) AS avg_rating
FROM reviews
GROUP BY product_id, TO_CHAR(review_date, 'YYYY-MM')
ORDER BY product_id, review_month;
```

**Sample Data:**

product_id	review_date	rating
101	2024-05-10	4
101	2024-05-15	5
101	2024-06-01	3



product_id	review_date	rating
102	2024-05-20	2
102	2024-06-05	4

Sample Output:

product_id	review_month	avg_rating
101	2024-05-01	4.5
101	2024-06-01	3.0
102	2024-05-01	2.0
102	2024-06-01	4.0

- **Note:** The `review_month` column format may vary by SQL dialect (e.g., `2024-05-01` or `2024-05`).

**Summary:** Use `GROUP BY` with a date truncation or formatting function to aggregate average ratings per product per month. Adjust the date function for your database system.

---